

Introduction

Motivation¹

The Great Music Service Company has a 1M catalog of songs to listen to available to their users. Given the large catalog of songs, no one user can be reasonably expected to go through all the content and choose what they want to listen to.

The Great Music Service conducted with survey of their active users and found that: (1) users are not aware of all the music available, (2) the catalog is too big and is hard to choose what to listen to, (3) they do not always know what they are looking for, and (4) they often use the music service as a song discovery tool in which they are explicitly interested in listening to new things.

Recommendations engines have shown to have significant benefit to commercial enterprises. Seventy-five percent of content watched in Netflix comes from its recommendation engine.² Moreover, thirty-five percent of Amazon's revenue comes from its recommendation system.³ Finally, studies show that 86% of consumer's purchases are significantly influenced by personalized recommendations.

Due to the survey results and recommender system statistics such as the ones highlighted above, the Chief Technology Officer and the Chief Data Scientist have decided to build a recommendation system for the music service platform to assist each user in finding and engaging with content that is relevant to them. If users engage with content relevant to them, they are more likely to spend more time in the service and less likely to churn. Higher engagement directly leads to higher revenue for the music service. The chief data scientist has tasked a member of the data science team to build a prototype recommendation system and recommend next steps.

Project Deliverables

¹ This scenario is used to illustrate a realistic use case for recommendation system and frame the project.

² <https://www.businessinsider.com/netflixs-recommendation-engine-drives-75-of-viewership-2012-4>

³ <http://rejoiner.com/resources/amazon-recommendations-secret-selling-online/>

What are recommendation systems?

Recommendation systems are a type of retrieval information technology. Recommendations engines differ from search in that: (1) users will not necessarily tell the engine what they are looking for, or (2) users may not know what they are looking for. These types of engine use individual's explicit preferences or behaviors (implicit preferences) to retrieve information that is relevant to the specific active user. The next section gives a brief overview of key recommender goals, types and evaluation metrics.

Brief Overview of Recommendation Systems

Goals of Recommendation Systems

- Prediction Version: Predict values a user will give to an unrated item.
- Ranking Version: Determine top-k items for a particular user.

Types of User Data Used in Recommendation Systems

- 1- **Explicit Rating:** Ratings are specified by the user in regards their level of preference. Like and dislike for the item is explicit. The task is to predict the rating for the unrated items in the matrix.
- 2- **Implicit Rating:** Tracking of user behavior. A missing value does not mean a dislike for an item. As such, it is recommended that the matrix is converted to a unary matrix (1,0) and missing values are filled with zeros to prevent overfitting.

Considerations

- ✓ Relevance
- ✓ Novelty
- ✓ Serendipity
- ✓ Diversity

Challenges

- 1- Sparsity
- 2- Cold Start
- 3- Fraud

Evaluation

- 1- MSE
- 2- RMSE
- 3- Precision
- 4- Recall
- 5- AUC
- 6- ROC
- 7- Kendall's Tau (for rankings)

Recommender Type	Description	Approaches	Advantages	Disadvantages
Content-Based	It recommends items that are similar in content to those the user has liked in the past or matches their 'profile'.	<ul style="list-style-type: none"> • Naïve Bayes classifier • K-nearest neighbor • Decision Trees • Neural networks 	<ul style="list-style-type: none"> • Easy to explain. • Provides personalized recommendations. 	<ul style="list-style-type: none"> • Leads to recommendations that are 'too' similar. • Does not have serendipity effects.
Collaborative Filtering	A user is recommended items based on past ratings of a group of users.	User-Based: the ratings provided by like-minded users of target A are used to make recommendations for A (i.e. neighborhood based)	<ul style="list-style-type: none"> • Easy to implement and explain. • Exploits the logic of 'word of mouth' and user similarities. 	<ul style="list-style-type: none"> • Does not scale well. • Computationally complex. • Highly engaged users can have a large impact. • Does not work well in sparse matrixes.
		Item-Based: First determine the set of items S that are most similar to target item B. The ratings of set S items are used to determine whether a user will like target item B.	<ul style="list-style-type: none"> • Easy to implement and explain. • Recommendations make sense. • Can be used as product association tool. • More efficient than user-based. 	<ul style="list-style-type: none"> • Can produce recommendations that are too 'obvious'. • Hard to discover items that are different, user-based tends to be better at this.
		Model Based: Decision trees, rule based, Bayesian & Latent Factor (matrix factorization, SVM, SVD). They work within an optimization framework.	<ul style="list-style-type: none"> • Tend to work reasonably well with sparse data. • Works within optimization framework. • Latent models assume similarity between users is induced by a 'unknown' factor in the data. 	<ul style="list-style-type: none"> • Need to regularize. • They are parametric models. • Can get stuck on local optimum cannot know if it has reached global optimum.
Hybrid	Combine several approaches in providing recommendations.	Usually is done by combining content-based and collaborative filtering methods through weighting, cascading or zipping.	<ul style="list-style-type: none"> • Exploit strengths of several approaches and minimize weaknesses of other approaches. 	<ul style="list-style-type: none"> • Complex to build and explain.

Knowledge Based	Recommendations based on explicit user requirements.	There are several approaches: conversational, search-based and navigation based.	<ul style="list-style-type: none"> • Works well when the user is knowledgeable about the item domain. • Works well for high-price ticket items. 	<ul style="list-style-type: none"> • Need a lot of information about the item.
Demographic Based	Recommendations based on user demographic attributes.	Use social graph and demographic features to build recommender.	<ul style="list-style-type: none"> • Uses demographic data. 	<ul style="list-style-type: none"> • Can lead to assumptions about certain demographic groups which may not be accurate for an individual.
Location Based	Take location into account when making recommendations.	Include location and geographical considerations into on other recommender approaches.	<ul style="list-style-type: none"> • Takes time into account when making recommendations potentially leading to higher relevance and/or accuracy. 	<ul style="list-style-type: none"> • Need location data.
Time Sensitive	Take time into account when making recommendations.	Include time and longitudinal considerations on other recommender approaches. Can be combined with location depending the domain.	<ul style="list-style-type: none"> • Takes time into account when making recommendations potentially leading to higher relevance and/or accuracy. 	<ul style="list-style-type: none"> • Needs time data. • It is longitudinal in nature. Need to be aware of time-effects.
Group Based	Recommendations to a group (i.e. find least miserable options for the group).	Instead of targeting recommendation to individuals, target group preferences.	<ul style="list-style-type: none"> • Makes recommendations for an entire group. • Attempts to find the most acceptable option for the group. 	<ul style="list-style-type: none"> • Works well in specific domains.

Data Wrangling

At a basic level, three pieces of information are needed to build a recommendation system: (1) information about the items, (2) information about item users, and (3) information about the relationship between items and users. As such, the goal of the data wrangling task was to obtain these three pieces of information. I have divided the data wrangling process into 12 parts. See figure 1 for a broad summary of the data wrangling steps.

Part 1. Data Extraction

File	Description	Source	Type	Quantity of Files
Million Song Dataset	Collection of 1M .h5 files. Each file represents data for one song.	Using the terminal: rsync -avzuP publicdata.opensciencedatacloud.org::ark:/31807/osdc-c1c763e4/remotefile /path/to/local_copy	.h5	1M
Million Song Summary File	Song metadata information.	http://labrosa.ee.columbia.edu/millionsong/sites/default/files/AdditionalFiles/msd_summary_file.h5	.h5	1
Artist Terms and Tags	Artist terms and tags obtained from the Echo Nest API and the Music Brainz.	Using the terminal: rsync -avzuP publicdata.opensciencedatacloud.org::ark:/31807/osdc-c1c763e4/remotefile /path/to/local_copy <i>**Located in file A</i>	SQLite	1
Million Song Duplicates	List of known song duplicates.	http://labrosa.ee.columbia.edu/millionsong/sites/default/files/AdditionalFiles/msd_duplicates.txt	.txt	1
Million Song Mismatches	List of known song mismatches.	http://labrosa.ee.columbia.edu/millionsong/sites/default/files/tasteprof/sid_mismatches.txt	.txt	1

Part 2. Song Metadata Summary File

The song metadata summary file is an .h5 file that has metadata information for 1M song files. The file is divided into three groups: (1) metadata, (2) analysis and (3) musicbrainz. Each group contains numpy arrays for each song.⁴

⁴ The aggregate song metafile was created by the researches of Columbia's LabRosa through iteration. The code they used can be found here:
https://github.com/tbertinmahieux/MSongsDB/blob/master/PythonSrc/create_aggregate_file.py

The **song metadata** group contains arrays that describe basic features of a song such as id, song title, song hotness, artist id, and release album. The **analysis metadata** group contains song's musical features such as key, length, danceability. The **musicbrainz** group includes the year of song release. The wrangling steps for each group of arrays are as follows:

- **Song Metadata Dataset:** The group's arrays were converted into a pandas dataframe. Once in the dataframe, object type columns were converted to strings and columns with b' ' string were removed. These strings were the result of python parsing the string data into Unicode from another format. There were also removed to make song identifying features consistent since these strings are not found in other complimentary datasets. The resulting data had missing values in some of the columns. The dataframe has 1M entries.
- **Analysis Metadata Dataset:** The group's arrays were converted into a pandas dataframe. Once in the dataframe, object type columns were converted to strings and columns with b' ' string were removed. These strings were the result of python parsing the string data into Unicode from another format. There were also removed to make song identifying features consistent. The resulting data had no missing values. The dataframe has 1M entries.
- **Musicbrainz:** The group's arrays were converted into a pandas dataframe. The dataset has two columns, with one of these columns representing year of song release. Zeroes in the year column represent missing. As such, zeroes were replaced with 'NaN' using `df.replace` and `np.nan`. The dataframe has 1M entries.

An alternative approach to using the song metadata summary file was to use an iterator to extract the information of interest from each of the individual song's .h5 files. The difference between the summary file and the full dataset, is that the full dataset contains artist similarity analysis and in-depth song analysis. This approach was explored; and a function, as well for loop were built to extract the data from the larger dataset. However, during testing the amount of time that it took to complete iterating through the files was substantial. Using the summary metadata file showed to be more efficient since the iteration work for the song's metadata has already been done. Moreover, for the purposes of building a recommendation system prototype, song metadata is sufficient and using in-depth song analysis or artist similarity is not required.

Part 3. Artist Terms and Tags

The artist terms and tags are contained in a SQLite database. The database has five tables. The tables `artist_mbttag` and `artist_term` were queried and converted to a pandas dataframe. The resulting dataframes were merged on artist id and using a left merge approach. Missing values in the merged dataframe were filled using the pandas method `fillna`.

Part 4. User Listening Data

The user listening data text file was converted to a dataframe. . The dataframe read the several columns of data as one. Separating by delimiter was used to separate data into several columns. The result was a dataframe with `user_id`, `song_id` and `play_count`. A column called “play” with value 1 was added to indicate the user played the song. The resulting dataframe has ~48.4M entries.

Part 5. Song Mismatches List

There is known song mismatches list. This means songs that were labeled incorrectly and matched incorrectly when the dataset was created. The song mismatches text file was converted to a dataframe. The dataframe read the several columns of data as one. Separating by delimiter was used to separate data into several columns. As a result, a dataframe with four columns was created: `song_id`, `track_id`, `song_name`, and `does_not_equal_to`.

Part 6. Song Duplicates List

There is known song duplicates list. The song duplicate text file was converted to a dataframe. The dataframe read the several columns of data as one. The structure of this data column is a song name followed by `track_ids` below.

The method `series.str.split` was used to split the `track_ids` and the song names. As a result of this a blank column was created, which was dropped. Two columns remained: `track_id` and song names. However, blanks were created in the `track_id` column because song names were previously there. These blanks were filled with NaN using `replace` method and `np.nan`. Further, the first five rows were dropped because they had dataset contact information and not actual duplicate information.

Moreover, the song names column has blanks cells below song names (i.e. they are parallel to `track_ids`). Using the pandas `fillna` with the forward fill method, song names were propagated

forward and to its respective track_ids. Track ids with blanks rows were dropped, extraneous number digits in front of the song names were removed, and the index was reset.

This list is the result of finding 131,661 items of 53,471 song objects. This means that 53,471 are unique values to the overall song dataset. To take this into account, the pandas duplicated method was used to identify unique and duplicated values. Finally, the dataframe was filtered to include only values that were found to be duplicates and remove the ones that are unique.

Part 7. Concatenating Song Metadata

The Song Metadata Dataset, Analysis Metadata Dataset and Musicbrainz datasets were concatenated on the index to create one large song metadata dataframe. The resulting dataframe has 1M entries. These columns have missing values: artist_familiarity, artist_hottnesss, artist_longitude, song_hottnesss and year.

Part 8. Removing Song Mismatches from Song Metadata

The concatenated song metadata file was merged with the song mismatches dataframe using a left merge approach and on song_id. Due to the merge new rows and columns were created, including a track_id_y which came from the song mismatches list. The resulting dataframe have ~1.0004M entries.

The track_id_y blanks were replaced with the string 'not mismatched'. This was done to indicate songs that did not merge with the mismatched items list. The dataframe was filtered by the 'not mismatched' string to remove mismatched items. The resulting dataframe has 981,022 entries.

The merge approach was used to filter duplicates rather than comparing the song_id series of each dataframe directly because each series' length and indexes are different making the comparison challenging and not obvious.

Part 9. Using Duplicates List to Remove Duplicates

The dataframe with mismatches removed was merged with the song duplicates dataframe using a left merge approach and on track_id. Due to the merge new rows and columns were created, including a song_name_y which came from the song duplicates. The resulting dataframe have 981,022 entries.

The song_name_y blanks were replaced with the string 'not duplicate'. This is to indicate songs that did not merged with the duplicates items list. The dataframe was filtered by the 'not duplicate' string to remove mismatched items. The resulting dataframe has 905,712 entries.

The merge approach was used to filter duplicates rather than comparing the track_id series of each dataframe directly because each series' length and indexes are different making the comparison challenging and not obvious.

Part 10. Removing columns created during merges and filling blanks

In this step, columns created by the merged process were dropped. As well blank entries were replaced with 'NaN' using two approaches: fillna and replace. The latter created two blank columns: analyzer version and genre. These blank columns were removed.

Part 11. Converting Datasets into Tidy Format⁵

As part of the wrangling process three main dataframes were created: (1) song metadata, (2) user listening, and (3) artist tags. Song metadata dataframe was melted using song_id as the identifying variable. User listening dataframe was melted using user_id as the identifying variable (the song_id is part of the values). Artist tags and terms dataframe was melted using artist_id as the identifying variable.

Part 12. Output Datasets into CSV files

The song metadata, user listening data and artist tags were exported to a local CSV file using chunksize 500K and the mode append to speed up exporting.

⁵ More info on tidy data here: <https://cran.r-project.org/web/packages/tidyr/vignettes/tidy-data.html>

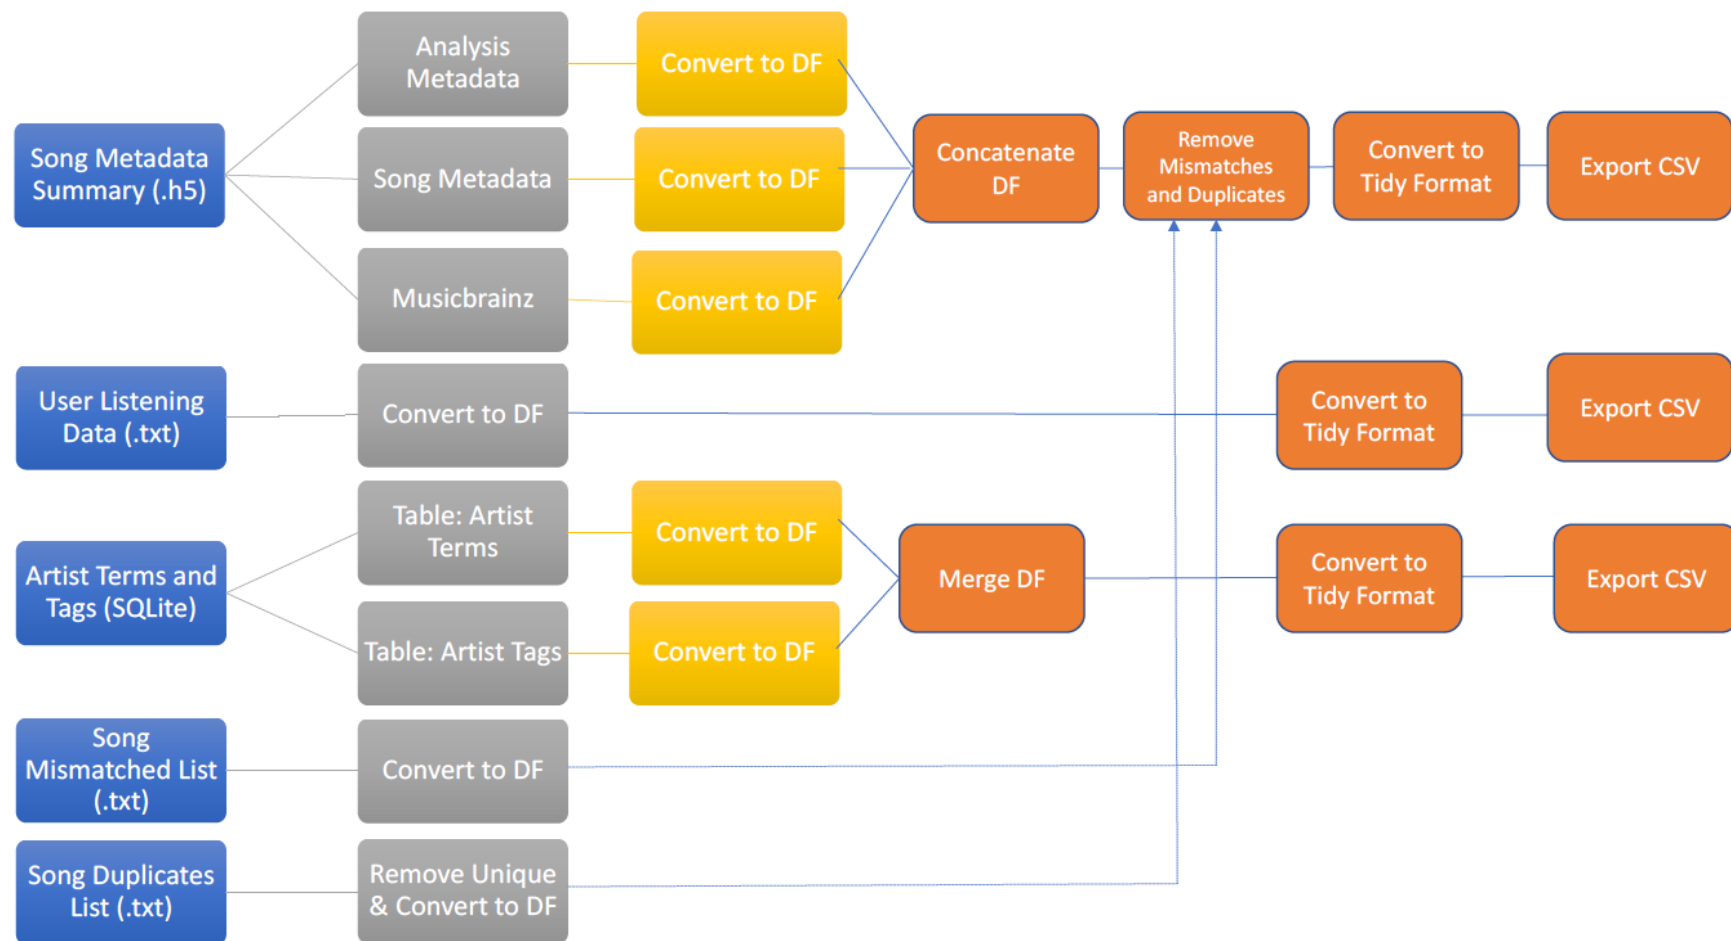


Figure 1. Summary of Data Wrangling Steps

Sources

- 1- Aggarwal, Charu C. Recommender Systems. Chapter 1 (pg 1-28) ISBN 978-3-319-29659-3
- 2- Asanov, Daniar. (2018). Algorithms and Methods in Recommender Systems.
https://www.snet.tu-berlin.de/fileadmin/fg220/courses/SS11/snet-project/recommender-systems_asanov.pdf
- 3- B. Akay, O. Kaynar and F. Demirkoparan, "Deep learning based recommender systems," *2017 International Conference on Computer Science and Engineering (UBMK)*, Antalya, 2017, pp. 645-648.
doi: 10.1109/UBMK.2017.8093489
- 4- Deep Learning based Recommender System: A Survey and New Perspectives. Zhang Shuai, Yao Lina, Sun Aixin. <https://arxiv.org/pdf/1707.07435.pdf>
- 5- Ellis, Daniel, Thierry Bertin-Mahieux. The Million Song Dataset Challenge.
<http://ecweb.ucsd.edu/~gert/papers/msdc.pdf>
- 6- Felfernig A., Jeran M., Ninaus G., Reinfrank F., Reiterer S., Stettinger M. (2014) Basic Approaches in Recommendation Systems. In: Robillard M., Maalej W., Walker R., Zimmermann T. (eds) Recommendation Systems in Software Engineering. Springer, Berlin, Heidelberg
- 7- <http://rejoiner.com/resources/amazon-recommendations-secret-selling-online/>
- 8- <https://www.businessinsider.com/netflixs-recommendation-engine-drives-75-of-viewership-2012-4>
- 9- Khusro, Shah & Ali, Zafar & Ullah, Irfan. (2016). Recommender Systems: Issues, Challenges, and Research Opportunities. 1179-1189. 10.1007/978-981-10-0557-2_112.
- 10- McFee, Brian, Lanckriet, Gret, Ellis, Daniel, Bertin-Mahieux. The Million Song Dataset Challenge. <http://ecweb.ucsd.edu/~gert/papers/msdc.pdf>
- 11- McFee, Brian, Lanckriet, Gret. LARGE-SCALE MUSIC SIMILARITY SEARCH WITH SPATIAL TREES. https://bmcftee.github.io/papers/ismir2011_sptree.pdf
- 12- Melville P., Sindhvani V. (2017) Recommender Systems. In: Sammut C., Webb G.I. (eds) Encyclopedia of Machine Learning and Data Mining. Springer, Boston, MA
- 13- Stanford Info Lab. Chapter 9: Recommendation Systems.
<http://infolab.stanford.edu/~ullman/mmds/ch9.pdf>