

RECOMMENDATION SYSTEM BUILD

Detailed Project Report

Ivette M Tapia

9.2.2018




Table of Contents

Introduction.....	3
Motivation	3
What are recommendation systems?	4
Brief Overview of Recommendation Systems	5
Dataset Creation Background	8
Data Wrangling.....	10
Dataset Features	18
Exploratory Data Analysis and Statistical Inference	19
Analytic Goals	19
Major Findings in Exploratory Data Analysis and Statistical Inference	19
Part 1. Song Features Exploration	21
Part 2. User Listen Counts	23
Part 3. User Listens Counts per Song.....	27
Part 4. User Listen Counts per Artist.....	31
Part 5. Relationships between basic song features and song listens	35
Part 6. Artist Tags and Terms	39

Introduction

This report summarizes the overall process taken to build a [item-item] or [matrix factorization] recommendation system algorithm. The data used to train and validate the recommendation system was the million song dataset and the user taste profile. This dataset was created and is distributed by Columbia University's LabRosa. The Million song dataset contains metadata for one million songs; while the user taste profile dataset contains user song listening counts and for 1.019M EchoNest users.

The report starts by outlining the motivation for the recommender system, and follows with a brief overview of recommendation systems. After this background, a broad overview of the process LabRosa took to create the datasets is outlined followed by the data wrangling steps taken to prepare the data for further analysis and recommender system creation. After this, exploratory data analysis results are discussed. Finally, the recommender system is described and predictive performance results are discussed. The coding language used throughout this project is python and the coding interface is Jupyter notebooks. The deliverables for the project is python code, a detailed report, and presentation slides.

Motivation¹

The Great Music Service company has a 1M catalog of songs to listen to available to their users. Given the large catalog of songs, no one user can be reasonably expected to go through all the content and choose what they want to listen to.

The Great Music Service conducted with survey of their active users found that: (1)²users are not aware of all the music available, (2) the catalog is too big and is hard to choose what to listen to, (3) they do not always know what they are looking for, and (4)

¹ This scenario is used to illustrate a realistic use case for recommendation system and frame the project.

² <https://www.businessinsider.com/netflixs-recommendation-engine-drives-75-of-viewership-2012-4>

they often use the music service as a song discovery tool in which they are explicitly interested in listening to new things.

Recommendations engines have shown to have significant benefit to commercial enterprises. Seventy-five percent of content watched in Netflix comes from its recommendation engine. Moreover, thirty-five percent of Amazon's revenue comes from its recommendation system.³ Finally, studies show that 86% of consumer's purchases are significantly influenced by personalized recommendations.

Due to the survey results and recommender system statistics, the Chief Technology Officer and the Chief Data Scientist have decided to build a recommendation system for the music service platform to assist each user in finding and engaging with content that is relevant to them. If users engage with content relevant to them, they are more likely to spend more time in the service and less likely to churn. Higher engagement directly leads to higher revenue for the music service. The chief data scientist has tasked a member of the data science team to build a prototype recommendation system and recommend next steps.

What are recommendation systems?

Recommendation systems are a type of retrieval information technology.

Recommendations engines differ from search in that: (1) users will not necessarily tell the engine what they are looking for, or (2) users may not know what they are looking for. These types of engine use individual's explicit preferences or behaviors (implicit preferences) to retrieve information that is relevant to the specific active user. The next section gives a brief overview of key recommender goals, types and evaluation metrics

³ <http://rejoiner.com/resources/amazon-recommendations-secret-selling-online/>

Brief Overview of Recommendation Systems

Goals of Recommendation Systems

- Prediction Version: Predict values a user will give to an unrated item.
- Ranking Version: Determine top-k items for a particular user.

Types of User Data Used in Recommendation Systems

- 1- **Explicit Rating:** Ratings are specified by the user in regards their level of preference. Like and dislike for the item is explicit. The task is to predict the rating for the unrated items in the matrix.
- 2- **Implicit Rating:** Tracking of user behavior. A missing value does not mean a dislike for an item. As such, it is recommended that the matrix is converted to a unary matrix (1,0) and missing values are filled with zeros to prevent overfitting.

Considerations

- ✓ Relevance
- ✓ Novelty
- ✓ Serendipity
- ✓ Diversity

Challenges

- 1- Sparsity
- 2- Cold Start
- 3- Fraud

Evaluation

- 1- MSE
- 2- RMSE
- 3- Precision
- 4- Recall
- 5- AUC
- 6- ROC
- 7- Kendall's Tau (for rankings)

Recommender Type	Description	Approaches	Advantages	Disadvantages
Content-Based	It recommends items that are similar in content to those the user has liked in the past or matches their 'profile'.	<ul style="list-style-type: none"> • Naïve Bayes classifier • K-nearest neighbor • Decision Trees • Neural networks 	<ul style="list-style-type: none"> • Easy to explain. • Provides personalized recommendations. 	<ul style="list-style-type: none"> • Leads to recommendations that are 'too' similar. • Does not have serendipity effects.
Collaborative Filtering	A user is recommended items based on past ratings of a group of users. These types of recommenders can be divided in two broad groups: memory based and model –based.	User-Based: the ratings provided by like-minded users of target A are used to make recommendations for A (i.e. neighborhood based)	<ul style="list-style-type: none"> • Easy to implement and explain. • Exploits the logic of 'word of mouth' and user similarities. 	<ul style="list-style-type: none"> • Does not scale well. • Computationally complex. • Highly engaged users can have a large impact. • Does not work well in sparse matrixes.
		Item-Based: First determine the set of items S that are most similar to target item B. The ratings of set S items are used to determine whether a user will like target item B.	<ul style="list-style-type: none"> • Easy to implement and explain. • Recommendations make sense. • Can be used as product association tool. • More efficient than user-based. 	<ul style="list-style-type: none"> • Can produce recommendations that are too 'obvious'. • Hard to discover items that are different, user-based tends to be better at this.
		Model Based: Decision trees, rule based, Bayesian & Latent Factor (matrix factorization, SVM, SVD). They work within an optimization framework.	Tend to work reasonably well with sparse data. Works within optimization framework. Latent models assume similarity between users is induced by a 'unknown' factor in the data.	Need to regularize. They are parametric models. Can get stuck on local optimum cannot know if it has reached global optimum.
Hybrid	Combine several approaches in providing recommendations.	Usually is done by combining content-based and collaborative filtering methods through weighting, cascading or zipping.	<ul style="list-style-type: none"> • Exploit strengths of several approaches and minimize weaknesses of other approaches. 	<ul style="list-style-type: none"> • Complex to build and explain.

Knowledge Based	Recommendations based on explicit user requirements.	There are several approaches: conversational, search-based and navigation based.	<ul style="list-style-type: none"> • Works well when the user is knowledgeable about the item domain. • Works well for high-price ticket items. 	<ul style="list-style-type: none"> • Need a lot of information about the item.
Demographic Based	Recommendations based on user demographic attributes.	Use social graph and demographic features to build recommender.	<ul style="list-style-type: none"> • Uses demographic data. 	<ul style="list-style-type: none"> • Can lead to assumptions about certain demographic groups which may not be accurate for an individual.
Location Based	Take location into account when making recommendations.	Include location and geographical considerations into on other recommender approaches.	<ul style="list-style-type: none"> • Takes time into account when making recommendations potentially leading to higher relevance and/or accuracy. 	<ul style="list-style-type: none"> • Need location data.
Time Sensitive	Take time into account when making recommendations.	Include time and longitudinal considerations on other recommender approaches. Can be combined with location depending the domain.	<ul style="list-style-type: none"> • Takes time into account when making recommendations potentially leading to higher relevance and/or accuracy. 	<ul style="list-style-type: none"> • Needs time data. • It is longitudinal in nature. Need to be aware of time-effects.
Group Based	Recommendations to a group (i.e. find least miserable options for the group).	Instead of targeting recommendation to individuals, target group preferences.	<ul style="list-style-type: none"> • Makes recommendations for an entire group. • Attempts to find the most acceptable option for the group. 	<ul style="list-style-type: none"> • Works well in specific domains.

Dataset Creation Background⁴

- The million-song dataset was created by using python wrapper on the ECHO Nest API for 10 days in December of 2010.
- The collection criteria for the songs dataset was as follows:
 1. Getting the most 'familiar' artists according to The Echo Nest, then downloading as many songs as possible from each of them.
 2. Getting the 200 top terms from The Echo Nest, then using each term as a descriptor to find 100 artists, then downloading as many of their songs as possible.
 3. Getting the songs and artists from the CAL500 dataset.
 4. Getting 'extreme' songs from The Echo Nest search params, e.g. songs with highest energy, lowest energy, tempo, song hottnesss, etc...
 5. A random walk along the similar artists links starting from the 100 most familiar artists.
- The resulting dataset is a collection of popular western songs. The number of songs was approximately 8,950 after step 1), step 3) added around 15,000 songs, and then added approx. 500,000 songs before starting step 5.
- The user dataset was created in October 2011 by the Echo Nest. The user data include a shuffled hash of persistent session identifiers from a small random selection of the musical universe and only play counts associated with Echo Nest song IDs that overlap with the song set. No usernames, listener details, original IDs, dates, IPs, locations were released.

⁴ <https://labrosa.ee.columbia.edu/millionsong/faq>, <http://blog.echonest.com/post/11992136676/taste-profiles-get-added-to-the-million-song>

- The user data is composed of solely of implicit feedback. This means that I only know what the user listened to, and a listen to a song does not mean the user liked or disliked the song.

Data Wrangling

At a basic level, three pieces of information are needed to build a recommendation system: (1) information about the items, (2) information about item users, and (3) information about the relationship between items and users. As such, the goal of the data wrangling task was to obtain these three pieces of information. I have divided the data wrangling process into 12 parts. See figure 1 for a broad summary of the data wrangling steps.

Part 1. Data Extraction

File	Description	Source	Type	Quantity of Files
Million Song Dataset	Collection of 1M .h5 files. Each file represents data for one song.	Using the terminal: rsync -avzuP publicdata.opensciencedatacloud.org::ark:/31807/osdc-c1c763e4/remotefile /path/to/local_copy	.h5	1M
Million Song Summary File	Song metadata information.	http://labrosa.ee.columbia.edu/millionsong/sites/default/files/AdditionalFiles/msd_summary_file.h5	.h5	1
Artist Terms and Tags	Artist terms and tags obtained from the Echo Nest API and the Music Brainz.	Using the terminal: rsync -avzuP publicdata.opensciencedatacloud.org::ark:/31807/osdc-c1c763e4/remotefile /path/to/local_copy **Located in file A	SQLite	1
Million Song Duplicates	List of known song duplicates.	http://labrosa.ee.columbia.edu/millionsong/sites/default/files/AdditionalFiles/msd_duplicates.txt	.txt	1
Million Song Mismatches	List of known song mismatches.	http://labrosa.ee.columbia.edu/millionsong/sites/default/files/tasteprofile/sid_mismatches.txt	.txt	1

Part 2. Song Metadata Summary File

The song metadata summary file is an .h5 file that has metadata information for 1M song files. The file is divided into three groups: (1) metadata, (2) analysis and (3) musicbrainz. Each group contains numpy arrays for each song.⁵

The **song metadata** group contains arrays that describe basic features of a song such as id, song title, song hotness, artist id, and release album. The **analysis metadata** group contains song's musical features such as key, length, danceability. The **musicbrainz** group includes the year of song release. The wrangling steps for each group of arrays are as follows:

- **Song Metadata Dataset:** The group's arrays were converted into a pandas dataframe. Once in the dataframe, object type columns were converted to strings and columns with b' ' string were removed. These strings were the result of python parsing the string data into Unicode from another format. There were also removed to make song identifying features consistent since these strings are not found in other complimentary datasets. The resulting data had missing values in some of the columns. The dataframe has 1M entries.
- **Analysis Metadata Dataset:** The group's arrays were converted into a pandas dataframe. Once in the dataframe, object type columns were converted to strings and columns with b' ' string were removed. These strings were the result of python parsing the string data into Unicode from another format. There were also removed to make song identifying features consistent. The resulting data had no missing values. The dataframe has 1M entries.

⁵ The aggregate song metadatafile was created by the researches of Columbia's LabRosa through iteration. The code they used can be found here:
https://github.com/tbertinmahieux/MSongsDB/blob/master/PythonSrc/create_aggregate_file.py

- **Musicbrainz:** The group's arrays were converted into a pandas dataframe. The dataset has two columns, with one of these columns representing year of song release. Zeroes in the year column represent missing. As such, zeroes were replaced with 'NaN' using `df.replace` and `np.nan`. The dataframe has 1M entries.

An alternative approach to using the song metadata summary file was to use an iterator to extract the information of interest from each of the individual song's .h5 files. The difference between the summary file and the full dataset, is that the full dataset contains artist similarity analysis and in-depth song analysis. This approach was explored; and a function, as well for loop were built to extract the data from the larger dataset.

However, during testing the amount of time that it took to complete iterating through the files was substantial. Using the summary metadata file showed to be more efficient since the iteration work for the song's metadata has already been done. Moreover, for the purposes of building a recommendation system prototype, song metadata is sufficient and using in-depth song analysis or artist similarity is not required.

Part 3. Artist Terms and Tags

The artist terms and tags are contained in a SQLite database. The database has five tables. The tables `artist_mbttag` and `artist_term` were queried and converted to a pandas dataframe. The resulting dataframes were merged on artist id and using a left merge approach. Missing values in the merged dataframe were filled using the pandas method `fillna`.

Part 4. User Listening Data

The user listening data text file was converted to a dataframe. . The dataframe read the several columns of data as one. Separating by delimiter was used to separate data into several columns. The result was a dataframe with `user_id`, `song_id` and `play_count`. A

column called “play” with value 1 was added to indicate the user played the song. The resulting dataframe has ~48.4M entries.

Part 5. Song Mismatches List

There is known song mismatches list. This means songs that were labeled incorrectly and matched incorrectly when the dataset was created. The song mismatches text file was converted to a dataframe. The dataframe read the several columns of data as one. Separating by delimiter was used to separate data into several columns. As a result, a dataframe with four columns was created: song_id, track_id, song_name, and does_not_equal_to.

Part 6. Song Duplicates List

There is known song duplicates list. The song duplicate text file was converted to a dataframe. The dataframe read the several columns of data as one. The structure of this data column is a song name followed by track_ids below. The method series.str.split was used to split the track_ids and the song names.

As a result of this a blank column was created, which was dropped. Two columns remained: track_id and song names. However, blanks were created in the track_id column because song names were previously there. These blanks were filled with NaN using replace method and np.nan. Further, the first five rows were dropped because they had dataset contact information and not actual duplicate information.

Moreover, the song names column has blanks cells below song names (i.e. they are parallel to track_ids). Using the pandas fillna with the forward fill method, song names were propagated forward and to its respective track_ids. Track ids with blank rows were dropped, extraneous number digits in front of the song names were removed, and the index was reset.

This list is the result of finding 131,661 items of 53,471 song objects. This means that 53,471 are unique values to duplicates list dataset. To take this into account, the pandas duplicated method was used to identify unique and duplicated values. Finally, the dataframe was filtered to include only values that were found to be duplicates and remove the ones that are unique.

Part 7. Concatenating Song Metadata

The Song Metadata Dataset, Analysis Metadata Dataset and Musicbrainz datasets were concatenated on the index to create one large song metadata dataframe. The resulting dataframe has 1M entries. These columns have missing values: artist_familiarity, artist_hottnesss, artist_longitude, song_hottnesss and year.

Part 8. Removing Song Mismatches from Song Metadata

The concatenated song metadata file was merged with the song mismatches dataframe using a left merge approach and on song_id. Due to the merge, new rows and columns were created, including a track_id_y which came from the song mismatches list. The resulting dataframe have ~1.0004M entries.

The track_id_y blanks were replaced with the string 'not mismatched'. This was done to indicate songs that did not merge with the mismatched items list. The dataframe was filtered by the 'not mismatched' string to remove mismatched items. The resulting dataframe has 981,022 entries.

The merge approach was used to filter duplicates rather than comparing the song_id series of each dataframe directly because each series' length and indexes are different making the comparison challenging and not obvious.

Part 9. Using Duplicates List to Remove Duplicates

The dataframe frame with mismatches removed was merged with the song duplicates dataframe using a left merge approach and on track_id. Due to the merge, new rows and columns were created, including a song_name_y which came from the song duplicates. The resulting dataframe have 981,022 entries.

The song_name_y blanks were replaced with the string 'not duplicate'. This is to indicate songs that did not merged with the duplicates items list. The dataframe was filtered by the 'not duplicate' string to remove mismatched items. The resulting dataframe has 905,712 entries.

The merge approach was used to filter duplicates rather than comparing the track_id series of each dataframe directly because each series' length and indexes are different making the comparison challenging and not obvious.

Part 10. Removing columns created during merges and filling blanks

In this step, columns created by the merged process were dropped. As well blank entries were replaced with 'NaN' using two approaches: fillna and replace. The latter created two blank columns: analyzer version and genre. These blank columns were removed.

Part 11. Converting Datasets into Tidy Format⁶

As part of the wrangling process three main dataframes were created: (1) song metadata, (2) user listening, and (3) artist tags. Song metadata dataframe was melted using song_id as the identifying variable. User listening dataframe was melted using user_id as the identifying variable (the song_id is part of the values). Artist tags and terms dataframe was melted using artist_id as the identifying variable.

⁶ More info on tidy data here: <https://cran.r-project.org/web/packages/tidyr/vignettes/tidy-data.html>

Part 12. Output Datasets into CSV files

The song metadata, user listening data and artist tags were exported to a local CSV file using chunksize 500K and the mode append to speed up exporting.

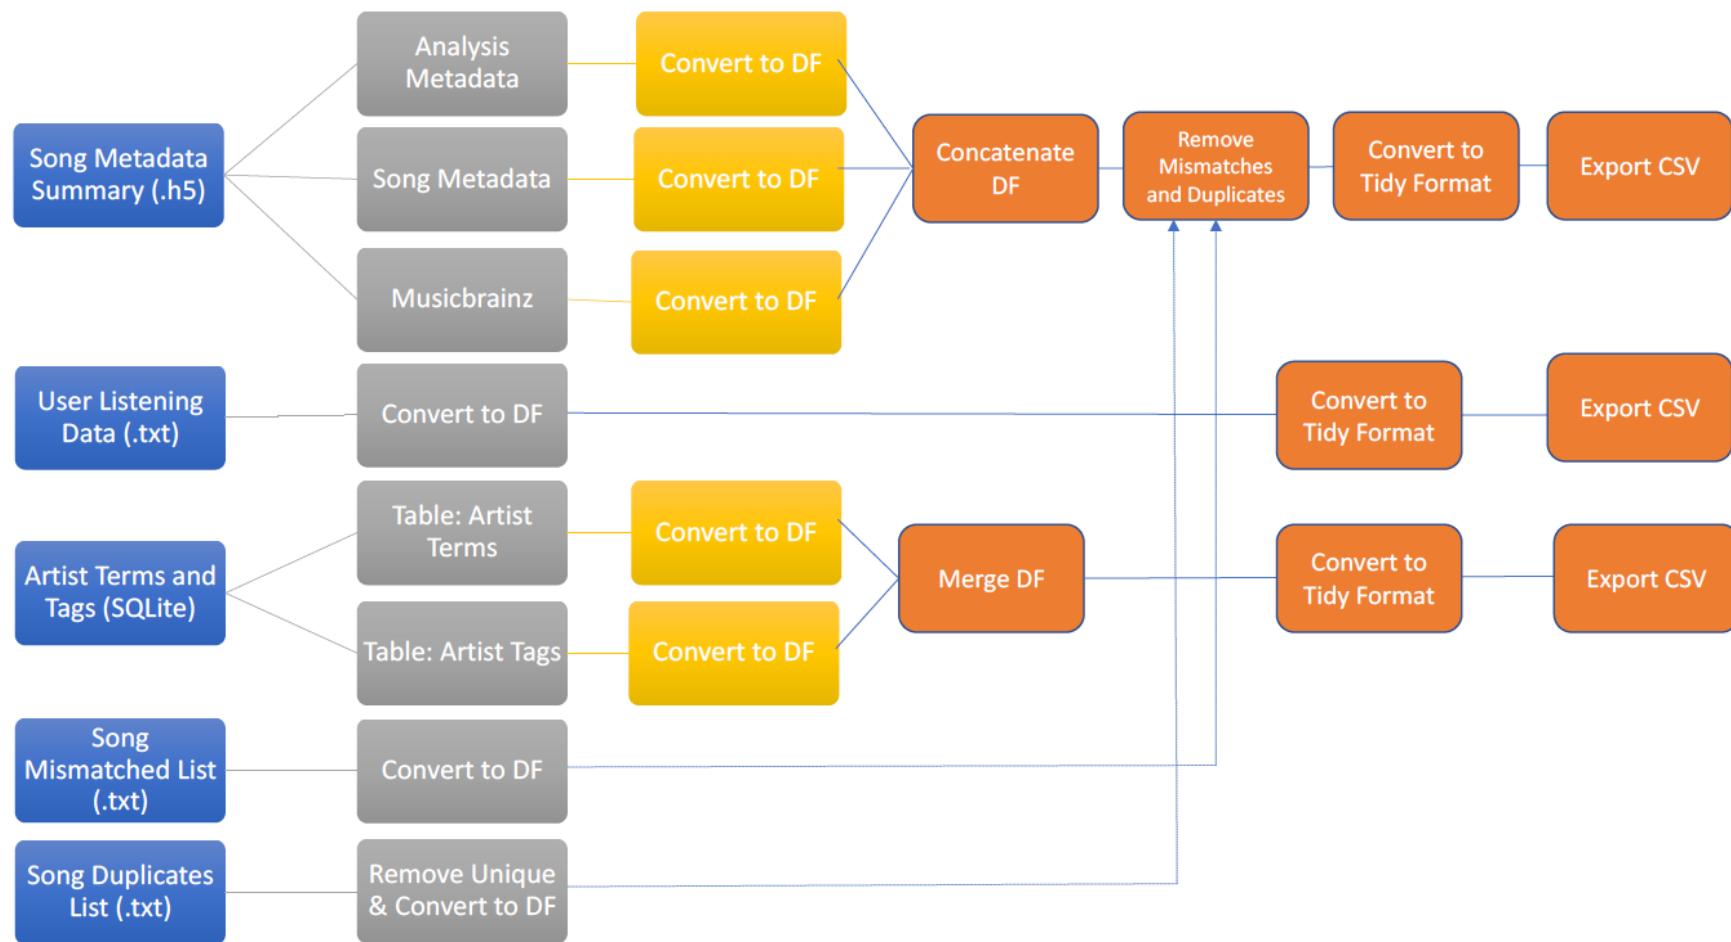


Figure 1. *Summary of Data Wrangling Steps*

Dataset Features⁷

Dataset	Feature Name	Description
User Listen Data	song_id	Echo Nest song ID
	user_id	Echo Nest user ID
	play_count	Number of times user played a song
	play	Whether a user played the song at least once. In recommender systems, it is best practice to treat implicit feedback as 0 or 1 (1 = consumed, 0 = not consumed)
Song Metadata	artist_7digitalid	ID from 7digital.com or -1
	artist_familiarity	algorithmic estimation
	artist_hottnesss	algorithmic estimation
	artist_id	Echo Nest ID artist ID
	artist_latitude	latitude
	artist_location	longitude
	artist_mbid	ID from musicbrainz.org
	artist_name	artist name
	artist_playmeid	ID from playme.com, or -1
	release	album name
	release_7digitalid	ID from 7digital.com or -1
	song_hottnesss	algorithmic estimation
	song_id	Echo Nest song ID
	title	song title
	track_7digitalid	ID from 7digital.com or -1
	audio_md5	audio hash code
	danceability	algorithmic estimation
	duration	in seconds
	energy	energy from listener point of view
	key	key the song is in
	loudness	overall loudness in dB
	mode	major or minor
	tempo	estimated tempo in BPM
	time_signature	estimate of number of beats per bar, e.g. 4
	track_id	Echo Nest track ID
	year	song release year from MusicBrainz or 0
Artists Tags	artist_id	Echo Nest artist id
	mbtag	Artists tags from musicbrainz.com
	terms	Artists terms from Echo Nest

⁷ <https://labrosa.ee.columbia.edu/millionsong/pages/example-track-description>

Exploratory Data Analysis and Statistical Inference

Analytic Goals

- 1- Calculate descriptive statistics of data points/features that will be used in the recommender system.
- 2- Get a sense of the contents of the song dataset.
- 3- Gain a strong sense of the distribution of the data points/features that will be used in the recommender system. In particular, user listen counts, song listen counts and artist listen counts.
- 4- Test hypotheses about the user listen counts, song listen counts and artist listen counts means.
- 5- Explore relationship between user interactions and basic song features.
- 6- Cursory view of the top artist tags and tags to gain a sense of top genres represented in the dataset.

Major Findings in Exploratory Data Analysis and Statistical Inference

Song Metadata	Descriptive Statistics	905,531 unique songs
	Year of release pattern	Peak year release songs in 2000's.
	Artist song counts pattern	Right skewed with 3 peaks.
Song indicator (y=1, n=0) aggregated by: user, song and artist	Descriptive Statistics	1.019M unique users 42% of song the catalog listened 66.5% of artists listened
	Normality Testing: <i>Graphical and Hypothesis testing</i>	Highly skewed to the right. Distributions are not normal . Some distributions may not satisfy central limit theorem.
	One - sample mean hypothesis testing	High likelihood sample means are close to population means. User Listen Indicator Means - User: 47 Song: 125 Artist: 1,367
Relationships between song features and play indicator	Joint plots with Pearson r	No strong relationship between listens and song features
Artist Tags	Top 20 artist tags from Echo Nest and <u>MusicBrainz</u>	Top tags and terms: Rock, electronic and pop.

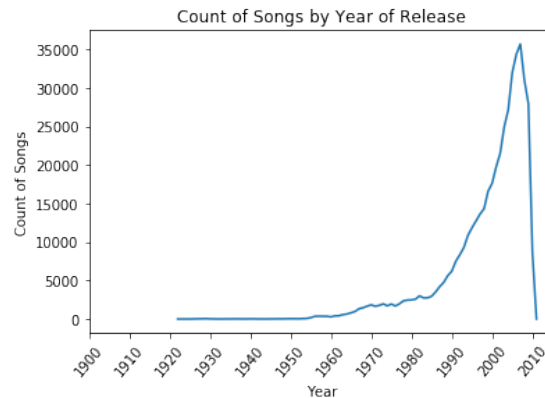
Part 1. Song Features Exploration

Basic Song Features Descriptive Statistics								
Song Feature	count	mean	std	min	25%	50%	75%	max
artist_familiarity	905,531	0.55	0.14	0.00	0.48	0.56	0.64	1.00
artist_hottnesss	905,700	0.38	0.12	0.00	0.33	0.38	0.44	1.08
song_hottnesss	551,532	0.35	0.23	0.00	0.22	0.38	0.53	1.00
danceability	905,712	0.00	0.00	0.00	0.00	0.00	0.00	0.00
duration	905,712	246.91	124.88	0.31	179.93	227.79	286.41	3034.91
energy	905,712	0.00	0.00	0.00	0.00	0.00	0.00	0.00
key	905,712	5.31	3.59	0.00	2.00	5.00	9.00	11.00
loudness	905,712	-10.16	5.24	-58.18	-12.74	-8.99	-6.39	4.32
mode	905,712	0.67	0.47	0.00	0.00	1.00	1.00	1.00
tempo	905,712	123.84	35.28	0.00	97.73	121.69	144.90	302.30
time_signature	905,712	3.59	1.23	0.00	3.00	4.00	4.00	7.00
year	456,811	1999	10.29	1922	1995	2002	2006	2011


– Finding #1: Song Features Descriptive Statistics

- Looking at the means it seems artist familiarity is 0.55, it falls somewhat in the middle of the possible values of 0 to 1.
- Artist hottness and song hottness mean values skew towards the lower range of hottness.
- In terms of song duration, the mean song is 4.10 minutes (which sounds accurate).
- The mean loudness of the tracks is -10.16, which suggests many of the tracks are not in the extreme of loudness (since this value is far away from max loudness of 4.31)
- The average key is 5.31 (this is around B major or D flat major).
- The mean beats per minute (tempo) is 123. Which means the mean song is in Allegro tempo: fast, quick, and bright.
- The mean mode is 0.66, there is an skew towards songs in the major scale.
- The mean time signature is 3.59 and the 75% percentile is 4. This means most songs are in 4/4 time signature.
- The mean year of release is 1998, with approximately. 25% of songs released between 2006 and 2011.

- **Finding #2:** There are 44,421 unique artists and 905,712 unique songs in the dataset.
- **Finding #3:** Year of release heavily skews to recent releases with a sharp peak in the 2000's. Year data is not complete, I have year data for 456,811 songs or 50% of the song dataset. It is unclear if the missing values are random or systematic, an example of this would be older songs tending to have missing values for year.



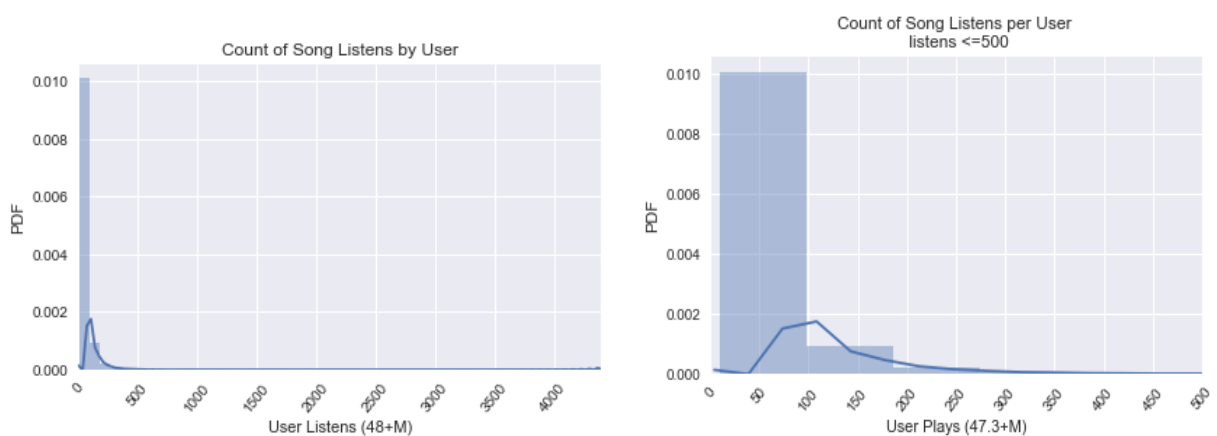
- **Finding #4:** Artist song counts are skewed to the right and have some peaks.



Part 2. User Listen Counts


Song Listens by User Descriptive Statistics								
	count	mean	std	min	25%	50%	75%	max
play	1,019,318	47.46	57.82	10.00	16.00	27.00	55.00	4,400
play_count	1,019,318	136.05	184.53	10.00	34.00	73.00	163.00	13,132

- **Finding #1:** There are 1,019,318 unique users.
- **Finding #2:** The play indicator has a mean of 47.45 and a standard deviation of 57.82. This means that the average user has listened to 47.45 different songs.
- **Finding #3:** The percentile distribution suggests a long tail somewhere. The 75% percentile is 55 songs and the max song plays is 4,400 songs.
- **Finding #4:** The distribution of user song counts is heavily skewed to the right. The kurtosis and skewedness statistics suggest high levels of positive skewedness and kurtosis.
 - Skewness = 4.826
 - Kurtosis = 68.503



– **Finding #5:** Distribution of user listen is not normally distributed.

- Graphical method shows that the user listen counts distribution vary significantly in shape from a theoretically normal distribution of the sample's mean and standard deviation.



- Hypothesis testing method shows that distribution is not normal. The null hypothesis (H_0) is that the distribution is normal. The alternative hypothesis (H_a) is that is not normal. The testing is at the alpha level of 0.05. Since in both tests the p-value is less than 0.05 the null hypothesis is rejected and the alternative hypothesis is accepted that the listens by user distribution is not normal.

- D'Agostino and Pearson's Normality Test⁸: stat=1059797.83, pvalue=0.0
- Anderson-Darling Normality Test⁹: stat=105,800.29 > 0.05 critical value of 0.787

– **Finding #6:** The user listen count data satisfies the assumptions of the central limit theorem.

- *Independence:* To satisfy this condition we will need to assume that the user's listen count is independent of the listen count of another user. This seems to me like a reasonable assumption. Nonetheless, it is unclear if this sample is 10% of the


⁸ This function tests the null hypothesis that a sample comes from a normal distribution. It is based on D'Agostino and Pearson's, test that combines skew and kurtosis to produce an omnibus test of normality.

⁹ The Anderson-Darling tests the null hypothesis that a sample is drawn from a population that follows a particular distribution. For the Anderson-Darling test, the critical values depend on which distribution is being tested against. This function works for normal, exponential, logistic, or Gumbel (Extreme Value Type I) distributions. This test has been shown to have more statistical power for testing normality.

population. The echo nest claims that this sample is a small subset of their music universe, how small is unclear. Find info

here: <http://blog.echonest.com/post/11992136676/taste-profiles-get-added-to-the-million-song>

- *Randomness*: The Echo Nest randomly selected a sample of users whose play counts matched to the song ID's in the dataset.
 - *Sample Size > 30*: The sample size is greater than 30.
- **Finding #8**: There is a 95% chance that the user listens population mean is between 47.34 - 47.57.
- Using a bootstrap approach with 10,000 trials, the confidence interval for the user listen counts is 47.34 - 47.57. The mean replicates are normally distributed (see figures below).
 - This confidence interval is narrow and contains the sample mean of 47.45.



- **Finding #9**: There is a significant probability that the population mean is 45.47 listens. The null hypothesis (H_0) is that the population mean is 45.47 user song listens. The alternative hypothesis (H_a) is that the population mean is not 45.47 user song listens. The alpha level is 0.05. Since in both tests the p-value is greater than 0.05, the null hypothesis that the population mean is 45.47 user song listens is cannot be rejected.
- Bootstrap hypothesis testing: pvalue = 0.4996
 - One sample t-test: stat=0.119, pvalue=0.905



User Listen Counts/Plays Conclusions

1. The distribution is not normal.
2. According to the bootstrap replicates of the mean and the t-test it is highly likely that the population user listen count mean is around 47.59 unique song listens per user.
3. There are several factors that could be driving these results and creating a heavily right skewed user listen count distribution:
 - The user play dataset is very large.
 - The data was collected in a certain timeframe and procedure and we do not have full user listening history.
 - The song data was extracted in December of 2011 which is holiday season, which could affect the user listen distribution (since only those who matched with the song dataset are included).
 - There are specific characteristics of users with high user-song interactions (such as business customer, several people in the same account etc.).
 - The structure of the music industry, in which a few songs/artists create hits and most working artists do not.


Part 3. User Listen Counts per Song

User Listens by Songs Descriptive Statistics								
	count	mean	std	min	25%	50%	75%	max
play	384,546	125.79	799.03	1.00	4.00	13.00	52.00	110,479
play_count	384,546	360.63	3,256.81	1.00	8.00	32.00	133.00	726,885

- **Finding #1:** There are 384,546 unique songs that have been listened to. This means that 42% of the song catalog has been listened to and 58% has not been listened to at all.
- **Finding #2:** As shown in the previous section, at most a user as listed to 4,400 unique songs, and the average user has listened to 47.45 different songs. We can conclude that the song catalog and user listens are sparsely distributed (most users have listened to a small portion of the available catalog).
- **Finding #3:** The play indicator has a mean of 125.79 and a standard deviation of 799. This means that the average song has 125.79 unique user listens.
- **Finding #4:** The percentile distribution suggests a long tail somewhere. The 75% percentile is 52 songs and the max song plays is 110,479 unique user listens.
 - Skewness = 46.11
 - Kurtosis = 3780.29



- **Finding #5:** Distribution of user listens per song is **not normally distributed**.



- Graphical method shows that the user song listens per track distribution vary significantly in shape from a theoretically normal distribution of the sample's mean and standard deviation.
- Hypothesis testing method shows that distribution is **not normal**. The null hypothesis (H_0) is that the distribution is normal. The alternative hypothesis (H_a) is that is not normal. The testing is at the alpha level of 0.05. Since in both tests the p-value is less than 0.05, the null hypothesis is rejected and the alternative hypothesis is accepted that the listens by user distribution is not normal.
 - D'Agostino and Pearson's Normality Test¹⁰: stat=1176273.18, pvalue=0.0
 - Anderson-Darling Normality Test¹¹: stat= 110035.21 > 0.05 critical value of 0.787

- **Finding #6:** It is suspect that the distribution satisfies the conditions of the central limit theorem.


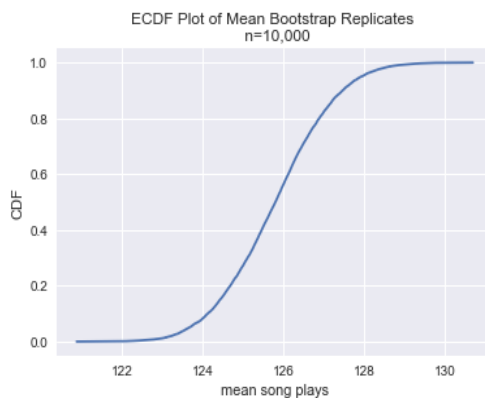
- *Independence:* To satisfy this condition we will need to assume that a song play count is independent of the play count of another song. In this case, I think this may not be not a reasonable assumption to make due to popularity, artists similarity,

¹⁰ This function tests the null hypothesis that a sample comes from a normal distribution. It is based on D'Agostino and Pearson's, test that combines skew and kurtosis to produce an omnibus test of normality.

¹¹ The Anderson-Darling tests the null hypothesis that a sample is drawn from a population that follows a particular distribution. For the Anderson-Darling test, the critical values depend on which distribution is being tested against. This function works for normal, exponential, logistic, or Gumbel (Extreme Value Type I) distributions. This test has been shown to have more statistical power for testing normality.

song similarity and how the echo nest service decided to show the catalog to their users. Also, it is unclear if this sample is 10% of the population.

- *Randomness*: The Echo Nest randomly selected a sample of users whose play counts matched to the song ID's in the dataset.
 - *Sample Size > 30*: The sample size is greater than 30.
- **Finding #8**: There is a 95% chance that the user listens population mean is between 47.34 - 47.57.
- Using a bootstrap approach with 10,000 trials, the confidence interval for the user listens per song is 123.31- 128.34. The mean replicates are normally distributed (see figures below).
 - This confidence interval contains the sample mean of 125.79.



- **Finding #9**: There is a significant probability that the population mean is 45.47 listens. The null hypothesis (H_0) is that the population mean is 45.47 user song listens. The alternative hypothesis (H_a) is that the population mean is not 45.47 user song listens. The alpha level of 0.05. Since in both tests the p-value is greater than 0.05, the null hypothesis that the population mean is 125.79 user song listens cannot be rejected. The t-test should be taken with some suspicion since this distribution may not satisfy the central limit theorem.
- Bootstrap hypothesis testing: pvalue = 0.496
 - One sample t-test: stat=0.003, pvalue= 0.997


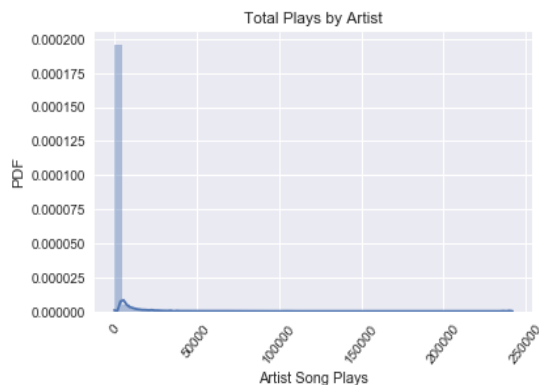
User Listens/Plays per Song Conclusions

- 1- The distribution is not normal.
- 2- According to the bootstrap replicates of the mean and the t-test it is highly likely that the population user listen count by song mean is around 125.79. The t-test should be taken with some suspicion since this distribution may not satisfy the central limit theorem.
- 3- There are several factors that could be driving these results and creating a heavily right skewed user listen count distribution:
 - a. The user play dataset is very large.
 - b. The data was in a certain timeframe and procedure and we do not have full user listening history.
 - c. The song data was extracted in December of 2011 which is holiday season, which could affect the user song play behavior.
 - d. There are specific characteristics of the user with high user play interactions (such as business customer, several people in the same account etc.).
 - e. The structure of the music industry, in which a few songs/artists create hits and most working artists do not.


Part 4. User Listen Counts per Artist

User Listens by Artist Descriptive Statistics								
	count	mean	std	min	25%	50%	75%	max
play	29,559.0	1,367.44	6,498.71	1.00	20.00	112.00	569.50	241,823.00
play_count	29,559.0	3,895.82	18,704.71	1.00	50.00	323.00	1,699.00	884,464.00

- **Finding #1:** There are 29,559 unique artists that have been listened to. This means that 66.5% of artists in the song dataset has been listened to, and 33.5% artists have not been listened to at all.
- **Finding #2:** The user listen indicator by artist has a mean of 1,367.44 and a standard deviation of 6,498.71. This means that the average artist has 1,367.44 unique user listens.
- **Finding #3:** The percentile distribution suggests a long tail somewhere. The 75% percentile is 569.50 songs and the max song plays is 241,823 unique user listens.



- **Finding #4:** Distribution of user listens per artist is not normally distributed.
 - Skewness = 15.13
 - Kurtosis = 344.37




- Graphical method shows that the user listens by artist distribution vary significantly in shape from a theoretically normal distribution of the sample's mean and standard deviation.
- Hypothesis testing method shows that distribution is **not normal**. The null hypothesis (H_0) is that the distribution is normal. The alternative hypothesis (H_a) is that is not normal. The testing is at the alpha level of 0.05. Since in both tests the p-value is less than 0.05, the null hypothesis is rejected and the alternative hypothesis is accepted that the listens by user distribution is not normal.
 - D'Agostino and Pearson's Normality Test¹²: stat= 56969.63, pvalue=0.0
 - Anderson-Darling Normality Test¹³: stat= 8001.32 > 0.05 critical value of 0.787
- **Finding #6:** It is suspect that the distribution satisfies the conditions of the central limit theorem.
 - *Independence:* To satisfy this condition we will need to assume that an artist play count is independent of the play count of another artist play count. In this case, I think this may not be not a reasonable assumption to make due to popularity effects, artists similarity, song similarity and how the echo nest service decided to show the catalog to their users. However, it may be a reasonable assumption that this sample represents <10% of the artist population.

¹² This function tests the null hypothesis that a sample comes from a normal distribution. It is based on D'Agostino and Pearson's, test that combines skew and kurtosis to produce an omnibus test of normality.

¹³ The Anderson-Darling tests the null hypothesis that a sample is drawn from a population that follows a particular distribution. For the Anderson-Darling test, the critical values depend on which distribution is being tested against. This function works for normal, exponential, logistic, or Gumbel (Extreme Value Type I) distributions. This test has been shown to have more statistical power for testing normality.

- *Randomness*: The Echo Nest randomly selected a sample of users whose play counts matched to the song ID's in the dataset.
 - *Sample Size > 30*: The sample size is greater than 30.
- **Finding #7**: There is a 95% chance that the artists listens population mean is between 1294.93- 1443.78.
- Using a bootstrap approach with 10,000 trials, the confidence interval for the user listens by artist is 1294.93- 1443.78. The mean replicates are normally distributed (see figures below).
 - This confidence interval contains the sample mean of 1,367.44.




- **Finding #9**: There is a significant probability that the population mean is 1,367.44 user listens per artist. The null hypothesis (H_0) is that the population mean is 1,367.44 user listens per artist. The alternative hypothesis (H_a) is that the population mean is not 1,367.44 user listens per artist. The alpha level of 0.05. Since in both tests the p-value is greater 0.05, the null hypothesis that the population mean is not 1,367.44 user listens per artists listens is cannot be rejected. The t-test should be taken with some suspicion since this distribution may not satisfy the central limit theorem.
- Bootstrap hypothesis testing: pvalue = 0.5033
 - One sample t-test: stat= -4.202, pvalue= 0.999

User Listens per Artist Conclusions

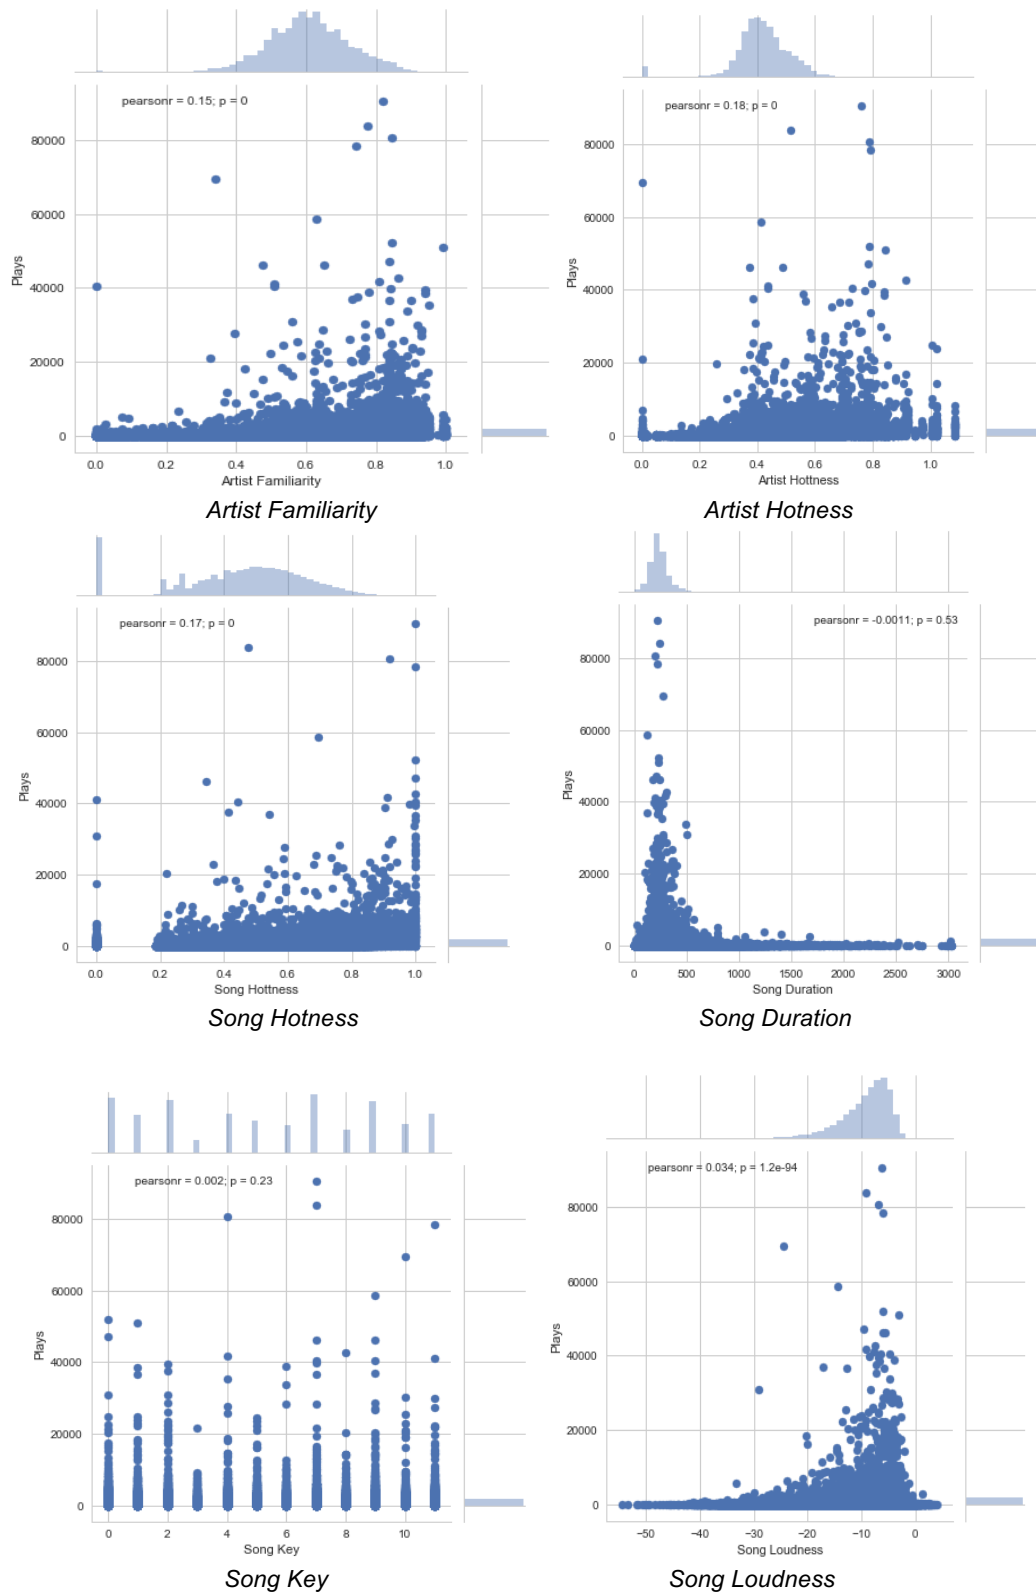
- 1- The distribution is not normal.
- 2- According to the bootstrap replicates of the mean and the t-test it is highly likely that the population user listen count per artist mean is around 1,367.44. The t-test should be taken with some suspicion since this distribution may not satisfy the central limit theorem.
- 3- There are several factors that could be driving these results and creating a heavily right skewed user listen per artist distribution:
 - a. The user play dataset is very large.
 - b. The data was in a certain timeframe and procedure and we do not have full user listening history.
 - c. The song data was extracted in December of 2011 which is holiday season, which could affect the user song play behavior.
 - d. There are specific characteristics of the user with high user play interactions (such as business customer, several people in the same account etc.).
 - e. The structure of the music industry, in which a few songs/artists create hits and most working artists do not.


Part 5. Relationships between basic song features and song listens

- There does not seem to be a strong relationship between unique song plays/interactions and any of the basic song features.
- The lack of relationship may be due to the sparsity of the data. For example, most users having not listened to that much of the catalog to result in meaningful patterns.
- The lack of relationship may also be due to the fact that many of the songs may be too heterogeneous or different.
- There may be patterns or correlations if features are looked at in combination rather than isolated.
- There seems to be a strong linear relationship between play count and plays. This means that there is a tendency in that the higher the play counts, the higher the number of unique user plays. This makes sense and ought to be expected given that the play indicator is derived from the play count.
- The high person correlation coefficient value of the relationship between user plays and user play counts leads me to think that song/artist popularity maybe an important factor in song/artist/user interactions (i.e. the more people hear about a song/artist the more likely they are to interact with it and have play it more).




Play Count and Plays







Song Mode




Song Tempo




Song Release Year



Time Signature



Part 6. Artist Tags and Terms



– MusicBrainz Tags

- There is a total of 2,321 unique MusicBrainz tags.
- The average tag has ~344 unique artists associated with it.
- There seems to be skewedness, 75% of the tags have 104 artists while the maximum artists associated with a tag is ~38.5K.
- As discussed previously there are 44,421 unique artists in the song dataset. So this means that 98% of artists in the MSD dataset has one or more MusicBrainz tag.
- Classic pop, rock, british, uk, american indie, english and folk are the most popular artists tags in the MusicBrainz data.

– EchoNest Terms

- There is a total of 7,643 unique terms.
- The average tag has ~212 unique artists associated with it.
- There seems to be some skewedness, 75% of the tags have 32 unique artists while the maximum artists associated with a tag is ~40.8K.
- As discussed previously, there are 44,421 unique artists in the song dataset. unique artists in the dataset. So this means that 98% of artists in the MSD dataset has one or more EchoNest tag.
- The Echo Nest has almost 3x more unique tags than the Music Brainz array.
- Rock, electronic, pop, hip hop, alternative and jazz the most popular artists terms in the Echo Nest data. Rock intersects with the top tags of the MusicBrainz.

Sources

- 1- Melville P., Sindhvani V. (2017) Recommender Systems. In: Sammut C., Webb G.I. (eds) Encyclopedia of Machine Learning and Data Mining. Springer, Boston, MA
- 2- Felfernig A., Jeran M., Ninaus G., Reinfrank F., Reiterer S., Stettinger M. (2014) Basic Approaches in Recommendation Systems. In: Robillard M., Maalej W., Walker R., Zimmermann T. (eds) Recommendation Systems in Software Engineering. Springer, Berlin, Heidelberg
- 3- Aggarwal, Charu C. Recommender Systems. Chapter 1 (pg 1-28) ISBN 978-3-319-29659-3
- 4- Stanford Info Lab. Chapter 9: Recommendation Systems.
<http://infolab.stanford.edu/~ullman/mmds/ch9.pdf>
- 5- McFee, Brian, Lanckriet, Gret. LARGE-SCALE MUSIC SIMILARITY SEARCH WITH SPATIAL TREES. https://bmcfree.github.io/papers/ismir2011_sptree.pdf
- 6- McFee, Brian, Lanckriet, Gret, Ellis, Daniel, Bertin-Mahieux. The Million Song Dataset Challenge. <http://ecweb.ucsd.edu/~gert/papers/msdc.pdf>
- 7- Ellis, Daniel, Thierry Bertin-Mahieux. The Million Song Dataset Challenge. <http://ecweb.ucsd.edu/~gert/papers/msdc.pdf>
- 8- Asanov, Daniar. (2018). Algorithms and Methods in Recommender Systems. https://www.snet.tu-berlin.de/fileadmin/fg220/courses/SS11/snet-project/recommender-systems_asanov.pdf
- 9- Khusro, Shah & Ali, Zafar & Ullah, Irfan. (2016). Recommender Systems: Issues, Challenges, and Research Opportunities. 1179-1189. 10.1007/978-981-10-0557-2_112.
- 10-Deep Learning based Recommender System: A Survey and New Perspectives. Zhang Shuai, Yao Lina, Sun Aixin. <https://arxiv.org/pdf/1707.07435.pdf>
- 11-B. Akay, O. Kaynar and F. Demirkoparan, "Deep learning based recommender systems," *2017 International Conference on Computer Science and Engineering (UBMK)*, Antalya, 2017, pp. 645-648.
doi: 10.1109/UBMK.2017.8093489

- 12-<https://www.businessinsider.com/netflixs-recommendation-engine-drives-75-of-viewership-2012-4>
- 13-<http://rejoiner.com/resources/amazon-recommendations-secret-selling-online/>
- 14-<https://labrosa.ee.columbia.edu/millionsong/pages/example-track-description>
- 15-<https://labrosa.ee.columbia.edu/millionsong/faq>
- 16-<http://blog.echonest.com/post/11992136676/taste-profiles-get-added-to-the-million-song>